

SIMATS ENGINEERING

ASSESSMENT TOOL 1

COURSE NAME: Artificial Intelligence

COURSE CODE: CSA1708

COURSE OUTCOME COVERED

CO1: Apply AI basics such as intelligent agents and uninformed search methods to solve problems effectively. (BL3)

Assessment Tool Weightage: 35%

QUESTIONS: 5

Duration : 1 Hour
Total Marks:50


Annexure A

Analytical Problem Solving

Code of Conduct:

*I **R.SACHIV(1921260172)** certify that this submission is my original work and that I have adhered to the guidelines specified for this assessment.
I understand that any violation of academic integrity rules will result in disciplinary action.*

Signature of the student:

A photograph of a piece of white paper with a handwritten signature in blue ink. The signature is stylized and appears to be 'R. SACHIV'.

Question 1: Water Jug Problem

Problem formulation

State: (x, y) where x = gallons in the 4-gallon jug ($0 \leq x \leq 4$) and y = gallons in the 3-gallon jug ($0 \leq y \leq 3$).

Initial state: $(0, 0)$

Goal state: $(2, n)$ – exactly 2 gallons in the 4-gallon jug (n can be anything).

Actions (production rules):

No.	Condition $\rightarrow$ Result	Meaning
1	$(x, y), x < 4 \rightarrow (4, y)$	Fill the 4-gallon jug from the pump
2	$(x, y), y < 3 \rightarrow (x, 3)$	Fill the 3-gallon jug from the pump
3	$(x, y), x > 0 \rightarrow (0, y)$	Empty the 4-gallon jug on the ground
4	$(x, y), y > 0 \rightarrow (x, 0)$	Empty the 3-gallon jug on the ground
5	$(x, y), x+y \geq 4, y > 0 \rightarrow (4, y-(4-x))$	Pour from 3-gal into 4-gal until the 4-gal is full
6	$(x, y), x+y \geq 3, x > 0 \rightarrow (x-(3-y), 3)$	Pour from 4-gal into 3-gal until the 3-gal is full
7	$(x, y), x+y \leq 4, y > 0 \rightarrow (x+y, 0)$	Pour all of the 3-gal into the 4-gal
8	$(x, y), x+y \leq 3, x > 0 \rightarrow (0, x+y)$	Pour all of the 4-gal into the 3-gal

Solution (found by breadth-first search – shortest, 6 steps)

Step	4-gal jug (x)	3-gal jug (y)	Action / Rule applied
0	0	0	Initial state
1	0	3	Fill the 3-gallon jug (Rule 2)
2	3	0	Pour all of the 3-gal into the 4-gal (Rule 7)
3	3	3	Fill the 3-gallon jug again (Rule 2)
4	4	2	Pour from the 3-gal into the 4-gal until it is full (Rule 5); 1 gallon goes in, 2 remain
5	0	2	Empty the 4-gallon jug (Rule 3)
6	2	0	Pour all of the 3-gal into the 4-gal (Rule 7) $\rightarrow$ GOAL

Alternative solution (also 6 steps): $(0,0) \rightarrow (4,0) \rightarrow (1,3) \rightarrow (1,0) \rightarrow (0,1) \rightarrow (4,1) \rightarrow (2,3)$. The last state has 2 gallons in the 4-gallon jug.

Result: Exactly 2 gallons is obtained in the 4-gallon jug; the state path is $(0,0) \rightarrow (0,3) \rightarrow (3,0) \rightarrow (3,3) \rightarrow (4,2) \rightarrow (0,2) \rightarrow (2,0)$.

Question 2: Mars Rover – PEAS Analysis

A Mars rover is an autonomous agent that travels on the Martian surface, and collects, analyses and explores soil/rock samples. Its design is analysed using the PEAS framework.

(a) Percepts (what the sensors give the agent)

- Camera images (stereo/panoramic, hazard-avoidance cameras) – terrain, rocks, slopes, obstacles, interesting samples.
- Spectrometer / chemical analysis readings (e.g. composition of rocks and soil, presence of water or minerals).

- Position and orientation: GPS-like navigation, wheel odometry, gyroscope, accelerometer, tilt sensors.
- Wheel/motor current and wheel slip – detects getting stuck in sand.
- Internal status: battery level, solar-panel power, temperature, fuel, system health.
- Environmental data: wind, dust, ambient temperature, radiation.
- Commands and messages from Earth control (with communication delay).

(b) Operating environment

Property	Characterisation
Observability	Partially observable – sensors have limited range; terrain beyond view is unknown.
Agents	Single agent (the rover) – other agents (Earth control) only send commands.
Determinism	Stochastic – wheel slip, dust storms, unpredictable terrain and sensor noise.
Episodic / Sequential	Sequential – a current movement affects future position, power and safety.
Static / Dynamic	Dynamic – weather, light, temperature and dust change while the rover is deciding.
Discrete / Continuous	Continuous – positions, speeds, time and sensor values are continuous.
Known / Unknown	Largely unknown – unexplored terrain; the rover must learn and map it.

(c) Actions (actuators)

- Move forward / backward, turn left / right, change speed, stop, climb or avoid obstacles (wheels/motors, steering).
- Deploy and move the robotic arm; drill, scoop or grind rock and soil; collect and store samples.
- Point cameras and instruments; take images; run spectrometer / on-board laboratory analysis.
- Orient solar panels, switch to sleep/power-saving mode, heat or cool components.
- Transmit data to Earth / orbiter and receive commands (antenna).

(d) Performance measure

- Number and scientific value of samples collected and analysed (quality and variety).
- Area of terrain explored and mapped; distance travelled towards target sites.
- Safety: number of collisions, falls, or times the rover got stuck (should be minimal).
- Energy efficiency: power used per task; the rover stays within its power budget.
- Mission lifetime and reliability (operating days without failure).
- Quantity and quality of data successfully transmitted to Earth; time taken per task.

(e) Most suitable agent architecture

A goal-based agent with utility-based decision making and learning (a hybrid / layered architecture) is the most suitable. Reasons:

- **Goals:** the rover has explicit goals (reach a site, analyse a rock) and must plan a sequence of actions – so a simple reflex agent is not enough.
- **Partial observability:** it must maintain an internal model/map of the world and its own state, so a model-based agent is needed.

- **Trade-offs:** there are competing aims (science value vs. battery vs. risk), so a utility function is needed to choose the best action.
- **Communication delay:** Earth cannot control it in real time, so the rover must act autonomously.
- **Learning:** the environment is unknown, so learning improves terrain classification and navigation over time.
- **Reflex layer:** a low-level reflex layer is added for immediate safety (stop if a cliff or tilt is detected).

Question 3: 8-Queens Problem Formulation

Place 8 queens on an 8×8 chessboard so that no queen attacks another (no two share a row, column or diagonal). In the figure given, the queen in the first column and the queen in the last column lie on the same diagonal, so it is not a solution (it has at least one attacking pair).

Problem components (incremental formulation)

Component	Definition
States	Any arrangement of 0 to 8 queens on the board such that no queen attacks another (queens are added column by column, leftmost first).
Initial state	The empty board (no queens placed).
Actions	Add a queen to a square in the leftmost empty column that is not attacked by any queen already on the board.
Transition model	Returns the board with the new queen added in the chosen square.
Goal test	8 queens are on the board and none is attacked.
Path cost	Not important – only the final configuration matters (each step has cost 1 if needed).

Search space comparison

Formulation	State space size	Remark
Complete-state: all 8 queens placed, any squares, move queens	$\approx 64 \times 63 \times \dots \times 57 \approx 1.8 \times 10^{14}$ sequences ($C(64,8) \approx 4.4 \times 10^9$ states)	Very large; uninformed search is impractical
Incremental, one queen per column, any row	$8^8 \approx 1.67 \times 10^7$	Smaller, but still contains attacking states
Incremental, add only non-attacked queens (chosen)	Only 2,057 states	Very small; easily solved by search

Exploring the search space with uninformed strategies

- **Depth-first search with backtracking:** place a queen in column 1, then in column 2 on a safe square, and so on. When a column has no safe square, backtrack to the previous column and try the next row. Memory needed is only $O(b \cdot m)$ and this is the usual method for 8-queens.
- **Breadth-first search:** generates all safe partial placements level by level (level $k = k$ queens). It finds all solutions but needs a lot of memory as the levels grow.
- **Iterative deepening / depth-limited search (limit = 8):** all solutions lie at depth exactly 8, so DFS with limit 8 is complete here without the memory cost of BFS.

Example solution (found by DFS with backtracking)

Row positions for columns 1 to 8: (1, 5, 8, 6, 3, 7, 2, 4). The 8-queens problem has 92 distinct solutions (12 unique up to rotation and reflection).

```
Q . . . . . . .
. . . . . Q .
. . . . Q . . .
. . . . . . Q
. Q . . . . .
. . . Q . . .
. . . . Q . .
. . Q . . . .
```

Check: all rows are different, one queen per column, and no two queens have the same value of (row - column) or (row + column), so there are no diagonal attacks.

Question 4: OLA Cab Booking – Problem-Solving Agent

Type of agent

A goal-based, problem-solving agent (with utility-based selection of the cab type) is suitable. The goal is clear (reach the destination from the source). The agent must look ahead and search over possible routes and cab choices, and select the option that best satisfies the customer's preferences: cost, comfort, time and availability. The environment is observable (map, traffic, fares and availability shown in the app), single-agent (the customer), mostly deterministic, sequential and dynamic (traffic, fares).

Problem formulation

Component	Definition
Initial state	Customer at pickup location with no ride booked.
Goal state	Customer reaches the destination location using the selected cab.
Actions	Enter pickup/drop; view cab types (Mini, Micro, Sedan, Shared, Prime...); select cab; confirm booking; track the driver; board and travel; pay.
Transition model	Each action moves the state: Searching → Cab selected → Booking confirmed → Driver assigned → Trip in progress → Destination reached.
Path cost / utility	Fare + waiting time + travel time (+ comfort preference). Lower cost and higher comfort give higher utility.

Pseudocode

```
function OLA-BOOKING-AGENT(percept) returns an action
  persistent: state, goal, problem, plan (sequence of actions), preferences

  state <- UPDATE-STATE(state, percept)           // location, traffic, fares
  if plan is empty then
    goal <- FORMULATE-GOAL(state)                 // reach D from S
    problem <- FORMULATE-PROBLEM(state, goal)     // states, actions, cost
    plan <- SEARCH(problem)                       // best route + cab
    if plan = failure then return NO-RIDE-AVAILABLE
  action <- FIRST(plan)
  plan <- REST(plan)
  return action

function SEARCH(problem) returns plan
  available <- GET-AVAILABLE-CABS(source)         // Mini, Micro, Sedan...
  best <- NULL; bestUtility <- -infinity
  for each cab in available do
    route <- FIND-SHORTEST-ROUTE(source, destination) // UCS or A* on map
    fare <- ESTIMATE-FARE(cab, route)
    eta <- ESTIMATE-TIME(cab, route)
    utility <- w1*comfort(cab) - w2*fare - w3*eta    // preference weights
    if utility > bestUtility then best <- (cab, route); bestUtility <- utility
```

```


if best = NULL then return failure
return [CONFIRM-BOOKING(best.cab), ASSIGN-DRIVER, TRACK-DRIVER,
        BOARD, TRAVEL(best.route), REACH-DESTINATION, PAY]

```

The agent then executes the plan step by step; if a percept changes (cab cancelled, traffic jam, fare change), the plan is empty or invalid and the agent re-plans.

Question 5: Uniform Cost Search – Warehouse Network

Graph (from the given figure)



Edges and costs: $S \rightarrow A = 1$, $S \rightarrow G = 12$, $A \rightarrow B = 3$, $A \rightarrow C = 1$, $B \rightarrow D = 3$, $C \rightarrow D = 1$, $C \rightarrow G = 2$, $D \rightarrow G = 3$.

Method

UCS always expands the node in the frontier (priority queue) with the lowest path cost $g(n)$. The goal test is applied when a node is selected for expansion, not when it is generated, which guarantees an optimal path.

Step-by-step execution

Step	Node expanded (g)	Nodes added / updated	Frontier after step (node:cost)
1	S (0)	$A = 0+1 = 1$; $G = 0+12 = 12$	A:1, G:12
2	A (1)	$B = 1+3 = 4$; $C = 1+1 = 2$	C:2, B:4, G:12
3	C (2)	$D = 2+1 = 3$; $G = 2+2 = 4$ (better than 12, so update)	D:3, B:4, G:4
4	D (3)	$G = 3+3 = 6$ (worse than 4, ignore)	B:4, G:4
5	B (4)	$D = 4+3 = 7$ (worse than 3, ignore)	G:4
6	G (4)	Goal selected for expansion → STOP	—

Explored (closed) list in order: S, A, C, D, B, G.

Result

Least-cost path: $S \rightarrow A \rightarrow C \rightarrow G$

Total minimum cost: $1 + 1 + 2 = 4$ units (fuel + time).

Interpretation: The direct route $S \rightarrow G$ costs 12, and the other route $S \rightarrow A \rightarrow C \rightarrow D \rightarrow G$ costs 6. The optimal route through warehouses A and C costs only 4, saving 8 units compared to the direct route and 2 units compared to the next best route. UCS is complete and optimal here because all edge costs are positive. Its time and space complexity is $O(b^{(1+\lceil C^*/\epsilon \rceil)})$, where C^* is the optimal cost and ϵ the minimum edge cost.

